


I-MAP: LOGBOOK (AI USAGE)

TSMK



LLM request

△ [ChatGPT](#) was used for this prompt.

My ros instance is inside a docker container, made using these files:

```
FROM ros:foxy

# Install dependencies
RUN apt update && apt install -y \
    python3-pip \
    git \
    && rm -rf /var/lib/apt/lists/*

# Setup workspace WORKDIR /root/microros_ws
# Clone micro-ROS
RUN git clone -b foxy https://github.com/micro-ROS/micro_ros_setup.git
src/micro_ros_setup

# Install dependencies
RUN . /opt/ros/foxy/setup.sh && \
    apt update && \
    rosdep update && \
    rosdep install --from-paths src --ignore-src -y

# Build workspace
RUN . /opt/ros/foxy/setup.sh && \
    colcon build

# Create agent workspace
RUN . /opt/ros/foxy/setup.sh && \
    . install/local_setup.sh && \
    ros2 run micro_ros_setup create_agent_ws.sh && \
    ros2 run micro_ros_setup build_agent.sh

# Copy entrypoint
COPY entrypoint.sh /entrypoint.sh
RUN chmod +x /entrypoint.sh
ENTRYPOINT ["/entrypoint.sh"]
```

```
#!/bin/bash

# Source ROS
source /opt/ros/foxy/setup.bash
source /root/microros_ws/install/local_setup.bash

# Run the agent
exec ros2 run micro_ros_agent micro_ros_agent serial --dev /dev/ttyACM0
```

```
version: "3.8"

services:
  ros2:
    build: .
    container_name: ros2_host
    network_mode: host
    privileged: true
    volumes:
      - /dev:/dev
    restart: unless-stopped
```

What i've seen from the micro-ros platformio page is that i can use a command like this:

```
docker run -it --rm -v /dev:/dev -v /dev/shm:/dev/shm --privileged --net=host  
microros/micro-ros-agent:$ROS_DISTRO serial --dev [YOUR BOARD PORT] -v6
```

What would be the best way to proceed?

LLM response

You're very close, but here's the key issue:

⇒ You're **mixing two different workflows**:

1. Building your own micro-ROS environment inside Docker
2. Using the **official prebuilt micro-ROS agent container**

That error (`rmw not found`) is happening because your custom image is **missing proper ROS 2 environment propagation during build/runtime**, and micro-ROS is particularly sensitive to that.

The core problem in your Dockerfile

This line:

```
RUN . /opt/ros/foxy/setup.sh && colcon build
```

✓ Works **only inside that RUN layer**

✗ But the environment is **lost in the next layers**

So later steps (and runtime) don't "see" `rmw`, even though ROS is installed.

✓ Best approach (recommended)

Option 1 — Use the official micro-ROS agent (simplest & most robust)

This is honestly the **best way unless you need customization**.

Just run:

```
docker run -it --rm \
  --privileged \
  --net=host \
  -v /dev:/dev \
  -v /dev/shm:/dev/shm \
  microros/micro-ros-agent:foxy \
  serial --dev /dev/ttyACM0 -v6
```

Why this is better:

- ✓ Already has correct `rmw` setup
- ✓ No build issues
- ✓ Maintained by micro-ROS team
- ✓ Matches PlatformIO expectations

⇒ If your goal is just:

"connect microcontroller → ROS 2" → use this

Option 2 — Fix your Dockerfile (if you want custom image)

If you *do* want your own container, here's what's wrong and how to fix it.

Problem 1: Environment not persistent

Fix by switching to SHELL:

```
SHELL ["/bin/bash", "-c"]
```

Then you can safely do:

```
RUN source /opt/ros/foxy/setup.bash && colcon build
```

Problem 2: You're building too much unnecessarily

You don't actually need:

```
colcon build
```

for `micro_ros_setup` itself.

Instead, the correct flow is:

```
RUN source /opt/ros/foxy/setup.bash && \
    ros2 run micro_ros_setup create_agent_ws.sh && \
    ros2 run micro_ros_setup build_agent.sh
```

Problem 3: Missing rosdep init (sometimes required)

Add:

```
RUN rosdep init || true
```

```
RUN rosdep update
```

Problem 4: rmw not found = ROS not sourced at runtime

Your entrypoint is actually good:

```
source /opt/ros/foxy/setup.bash
```

```
source /root/microros_ws/install/local_setup.bash
```

👉 That part is correct.

Minimal fixed Dockerfile

Here's a cleaner version:

```
FROM ros:foxy
```

```
SHELL ["/bin/bash", "-c"]
```

```
RUN apt update && apt install -y \
    python3-pip \
    git \
    && rm -rf /var/lib/apt/lists/*
```

```
WORKDIR /root/microros_ws
```

```
RUN git clone -b foxy https://github.com/micro-ROS/micro_ros_setup.git
src/micro_ros_setup
```

```
RUN source /opt/ros/foxy/setup.bash && \
    apt update && \
    rosdep init || true && \
    rosdep update && \
    rosdep install --from-paths src --ignore-src -y
```

```
RUN source /opt/ros/foxy/setup.bash && \
    ros2 run micro_ros_setup create_agent_ws.sh && \
```

```
ros2 run micro_ros_setup build_agent.sh
```

```
COPY entrypoint.sh /entrypoint.sh
```

```
RUN chmod +x /entrypoint.sh
```

```
ENTRYPOINT ["/entrypoint.sh"]
```

⚖️ Which should YOU choose?

👉 Use official container if:

- You just want micro-ROS agent running
- You're using PlatformIO
- You don't need custom ROS packages

👉 Use custom Dockerfile if:

- You want to integrate with other ROS 2 nodes
- You need custom builds or dependencies
- You're developing micro-ROS itself

💡 My honest recommendation

Given your setup (PlatformIO + micro-ROS):

👉 **Don't build your own image. Use the official one.**

Then, if needed, run your own ROS 2 container separately.

🦋 Bonus: docker-compose version (clean setup)

```
version: "3.8"
```

```
services:
```

```
  micro_ros_agent:
    image: microros/micro-ros-agent:foxy
    command: serial --dev /dev/ttyACM0 -v6
    network_mode: host
    privileged: true
    volumes:
      - /dev:/dev
    restart: unless-stopped
```

If you want, I can help you set up:

- 🔄 PlatformIO ↔ micro-ROS agent workflow
- 🐛 Debugging serial connection issues
- 🛠️ DDS/RMW tuning (FastDDS vs Cyclone)

Just tell me 🙋

LLM request

Say i use option 1 how can i put micro-ros on my teensy41 using platformio

LLM response

Great — this is exactly the setup micro-ROS + PlatformIO was designed for 👍

Using **Option 1 (official micro-ROS agent container)** is the cleanest path.
Let's walk through getting **micro-ROS running on your Teensy 4.1** step by step.

Big picture (so it all makes sense)

You'll have 2 sides:

■ **Microcontroller (Teensy 4.1)**

- Runs micro-ROS client (via PlatformIO)
- Sends data over **serial (USB)**

■ **Host (your PC / Docker)**

- Runs **micro-ROS agent**
- Bridges microcontroller ↔ ROS 2 DDS (rmw layer)

✓ **Step 1 — PlatformIO project setup**

1. Create a new PlatformIO project

Board:

teensy41

2. Add micro-ROS library

Edit your platformio.ini:

```
[env:teensy41]
```

```
platform = teensy
```

```
board = teensy41
```

```
framework = arduino
```

```
lib_deps =
```

```
    micro-ROS/micro_ros_platformio
```

```
build_flags =
```

```
    -D MICROROS_TRANSPORT_ARDUINO_SERIAL
```

Important: Configure micro-ROS transport

Create this file:

platformio.ini

extra_script.py

extra_script.py

```
Import("env")
```

```
env.Append(
```

```
    CPPDEFINES=[
```

```
        ("MICROROS_TRANSPORT_ARDUINO_SERIAL", 1)
```

```
    ]
```

```
)
```


✓ Step 2 — Minimal micro-ROS example

Replace your `src/main.cpp` with:

```
#include <micro_ros_platformio.h>
#include <rcl/rcl.h>
#include <rcl/rclc.h>
#include <std_msgs/msg/int32.h>

rcl_publisher_t publisher;
rcl_node_t node;
rclc_support_t support;
rcl_allocator_t allocator;

std_msgs__msg__Int32 msg;

void setup() {
  Serial.begin(115200);
  delay(2000);

  set_microros_serial_transports(Serial);

  allocator = rcl_get_default_allocator();

  rclc_support_init(&support, 0, NULL, &allocator);

  rclc_node_init_default(&node, "teensy_node", "", &support);

  rclc_publisher_init_default(
    &publisher,
    &node,
    ROSIDL_GET_MSG_TYPE_SUPPORT(std_msgs, msg, Int32),
    "teensy_counter");

  msg.data = 0;
}

void loop() {
  msg.data++;
  rcl_publish(&publisher, &msg, NULL);
  delay(1000);
}
```

✓ Step 3 — Build & upload

Connect your Teensy and run:

```
pio run -t upload
```

✓ Step 4 — Run micro-ROS agent

(Docker)

Now run the official agent:

```
docker run -it --rm \
  --privileged \
  --net=host \
  -v /dev:/dev \
  microros/micro-ros-agent:foxy \
  serial --dev /dev/ttyACM0 -v6
```

⇒ Replace /dev/ttyACM0 if needed:

```
ls /dev/ttyACM*
```

✓ Step 5 — Verify in ROS 2

Open another terminal (with ROS 2 Foxy sourced):

```
source /opt/ros/foxy/setup.bash
```

Check topics:

```
ros2 topic list
```

You should see:

```
/teensy_counter
```

Echo it:

```
ros2 topic echo /teensy_counter
```

🔥 Common issues (Teensy-specific)

✗ Device not found

Check:

```
ls /dev/ttyACM*
```

If missing:

- Replug USB
- Check permissions:

```
sudo chmod 666 /dev/ttyACM0
```

✗ Agent connects but no data

Usually means:

- Wrong baud rate (must match `Serial.begin()`)
- Transport not set correctly

✗ Nothing shows in ROS 2

Make sure:

- Agent is running
- Same `ROS_DOMAIN_ID` (if set)

💡 Tips for Teensy 4.1

- Use **USB (Serial)** → fastest + easiest
- Teensy is powerful → you can run timers, multiple publishers
- Avoid `delay()` in real projects → use timers (`rclcpp_executor`)

Optional next steps

Once this works, you can:

- Add subscribers (control motors, LEDs)
- Use sensors (IMU, encoder)
- Move to WiFi transport (ESP32-style)

Summary

➡ Best workflow for you:

- ✓ PlatformIO → micro-ROS firmware on Teensy
- ✓ Official Docker → micro-ROS agent
- ✓ ROS 2 on host → visualization/control

If you want, I can help you next with:

- 📡 Adding a subscriber (e.g. control LED from ROS 2)
- ⌚ Using timers instead of delay
- Debugging if your Teensy doesn't show up

Just tell me where you're stuck 🙌

LLM request

From step 3 it goes awry.

- 1) I get this error with your lib: `UnknownPackageError: Could not find the package with 'micro-ROS/micro_ros_platformio' requirements`

The official git page (https://github.com/micro-ROS/micro_ros_platformio) states to use:

```
lib_deps = https://github.com/micro-ROS/micro_ros_platformio
```

But this then in turn gives this error when run:

```
--- stderr: rmw_test_fixture
CMake Error at CMakeLists.txt:6 (find_package):
  By not providing "Findrmw.cmake" in
  CMAKE_MODULE_PATH this project has
  asked CMake to find a package configuration file provided by "rmw", but
  CMake did not find one.

Could not find a package configuration file provided by "rmw" with any of
the following names:

  rmwConfig.cmake
  rmw-config.cmake

Add the installation prefix of "rmw" to
CMAKE_PREFIX_PATH or set "rmw_DIR"
to a directory containing one of the above files. If "rmw" provides a
separate development package or SDK, be sure it has been installed.
```

Which i wanted to solve using the docker container (if possible, i don't know).

- 2) The docker now used uses `ros:foxy` is not supported (said on the same git page): `board_microros_distro = <distribution>`
Options:
 - humble
 - jazzy
 - kilted (default value)
 - rolling

But, including foxy, these all give errors when run.

How should i proceed?

LLM response

Yeah — what you're running into here is a **known pain point** with micro-ROS + PlatformIO right now:

⇒ **the library is not self-contained and expects a proper ROS 2 build environment during compilation**

Let's straighten this out cleanly, because right now you've got **two separate issues colliding**:

Root causes (what's actually breaking)

✗ 1. rmw error during PlatformIO build

This is the big one:

Could not find rmwConfig.cmake

⇒ This means:

- PlatformIO is trying to **build micro-ROS packages locally**
- But it **does NOT** have a **ROS 2 environment (rmw, rcl, etc.) available**

🔗 Important:

micro_ros_platformio is **not just a library** — it actually runs a mini ROS 2 build system (colcon) internally.

So:

- Your Docker container is irrelevant here ✗
- The build is happening on your **host machine via PlatformIO**

✗ 2. Wrong way of including the library

This:

lib_deps = https://github.com/micro-ROS/micro_ros_platformio

✓ Correct (per repo)

✗ But incomplete without extra config

✗ 3. ROS 2 Foxy is deprecated

You saw this correctly 🐞

micro-ROS now supports:

- humble
- jazzy
- kilted (default)
- rolling

⇒ **Foxy is effectively dead here**

✓ The clean solution (this actually works)

You have **two valid paths** — I'll give you the one that *won't fight you*.

🚀 ✓ **OPTION A (RECOMMENDED):**

Use Docker for BOTH build + agent

⇒ This avoids ALL rmw issues.

Instead of:

- PlatformIO building locally ✗

You:

- Build firmware **inside a ROS-enabled container** ✓

Step 1 — Use official micro-ROS dev container

Run:

```
docker run -it --rm \
  --privileged \
  -v $(pwd):/workspace \
  -v /dev:/dev \
  microros/micro-ros-agent:humble
```

Then inside container:

```
apt update
apt install -y python3-pip git
pip install platformio
```

Step 2 — Build inside container

```
cd /workspace
pio run -t upload
```

⇒ Now:

- ROS 2 is present ✓
- rmw exists ✓
- build succeeds ✓

OPTION B (harder): Fix local build

If you insist on building outside Docker:

✓ 1. Install ROS 2 Humble on host

⇒ This is REQUIRED (no workaround)

✓ 2. Source it BEFORE running PlatformIO

```
source /opt/ros/humble/setup.bash
pio run
```

✓ 3. Fix platformio.ini

```
[env:teensy41]
platform = teensy
board = teensy41
framework = arduino

lib_deps =
  https://github.com/micro-ROS/micro_ros_platformio

board_microros_distro = humble
```

✓ 4. Install missing rmw manually (just in case)

```
sudo apt install ros-humble-rmw-cyclonedds-cpp
```

⚠ Why your current setup fails

- Your Docker container:
 - o ✓ builds agent
 - o ✗ NOT used by PlatformIO build
- PlatformIO:
 - o ✗ runs outside ROS
 - o ✗ cannot find rmw

⇒ These are completely separate worlds right now

💡 Key insight (this unlocks everything)

micro-ROS PlatformIO = **cross-compiling ROS 2 packages**

So it **needs ROS 2 present at compile time**, not just runtime.

What I would do (fastest path)

1. ✗ Stop using your custom ROS Docker for now
2. ✓ Use official agent container (you already have this working idea)
3. ✓ Install ROS 2 Humble locally OR build inside Docker
4. ✓ Set:
5. `board_microros_distro = humble`

🚀 Minimal working combo

- Teensy firmware → built in ROS-enabled env
- Agent → Docker (microros/micro-ros-agent:humble)
- ROS 2 → same distro (humble)

If it STILL fails

Tell me:

- Your OS (Linux? WSL? Mac?)
- Full `platformio.ini`
- Exact error after setting humble

...and I'll debug it down to the exact missing package or env variable 🙌 Yeah — what you're running into here is a **known pain point** with micro-ROS + PlatformIO right now:

⇒ **the library is not self-contained and expects a proper ROS 2 build environment during compilation**

Let's straighten this out cleanly, because right now you've got **two separate issues colliding**:

Root causes (what's actually breaking)

✗ 1. rmw error during PlatformIO build

This is the big one:

Could not find rmwConfig.cmake

⇒ This means:

- PlatformIO is trying to **build micro-ROS packages locally**
- But it **does NOT have a ROS 2 environment (rmw, rcl, etc.) available**

🔗 Important:

micro_ros_platformio is **not just a library** — it actually runs a mini ROS 2 build system (colcon) internally.

So:

- Your Docker container is irrelevant here ✗
- The build is happening on your **host machine via PlatformIO**

✗ 2. Wrong way of including the library

This:

lib_deps = https://github.com/micro-ROS/micro_ros_platformio

✓ Correct (per repo)

✗ But incomplete without extra config

✗ 3. ROS 2 Foxy is deprecated

You saw this correctly 📖

micro-ROS now supports:

- humble
- jazzy
- kilted (default)
- rolling

⇒ Foxy is effectively dead here

✓ The clean solution (this actually works)

You have **two valid paths** — I'll give you the one that *won't fight you*.

👉 ✓ OPTION A (RECOMMENDED):

Use Docker for BOTH build + agent

⇒ This avoids ALL rmw issues.

Instead of:

- PlatformIO building locally ✗

You:

- Build firmware **inside a ROS-enabled container** ✓

Step 1 — Use official micro-ROS dev container

Run:


```
docker run -it --rm \
  --privileged \
  -v $(pwd):/workspace \
  -v /dev:/dev \
  microros/micro-ros-agent:humble
```

Then inside container:

```
apt update
apt install -y python3-pip git
pip install platformio
```

Step 2 — Build inside container

```
cd /workspace
pio run -t upload
```

☞ Now:

- ROS 2 is present ✓
- rmw exists ✓
- build succeeds ✓

OPTION B (harder): Fix local build

If you insist on building outside Docker:

✓ 1. Install ROS 2 Humble on host

☞ This is REQUIRED (no workaround)

✓ 2. Source it BEFORE running PlatformIO

```
source /opt/ros/humble/setup.bash
pio run
```

✓ 3. Fix platformio.ini

```
[env:teensy41]
platform = teensy
board = teensy41
framework = arduino

lib_deps =
  https://github.com/micro-ROS/micro_ros_platformio
```

```
board_microros_distro = humble
```

✓ 4. Install missing rmw manually (just in case)

```
sudo apt install ros-humble-rmw-cyclonedds-cpp
```

⚠ Why your current setup fails

- Your Docker container:
 - o ✓ builds agent
 - o ✗ NOT used by PlatformIO build
- PlatformIO:
 - o ✗ runs outside ROS
 - o ✗ cannot find rmw

👉 These are completely separate worlds right now

💡 Key insight (this unlocks everything)

micro-ROS PlatformIO = **cross-compiling ROS 2 packages**

So it **needs ROS 2 present at compile time**, not just runtime.

What I would do (fastest path)

1. ✗ Stop using your custom ROS Docker for now
2. ✓ Use official agent container (you already have this working idea)
3. ✓ Install ROS 2 Humble locally OR build inside Docker
4. ✓ Set:
5. `board_microros_distro = humble`

🦋 Minimal working combo

- Teensy firmware → built in ROS-enabled env
- Agent → Docker (microros/micro-ros-agent:humble)
- ROS 2 → same distro (humble)

If it STILL fails

Tell me:

- Your OS (Linux? WSL? Mac?)
- Full `platformio.ini`
- Exact error after setting humble

...and I'll debug it down to the exact missing package or env variable 🙌

LLM request

Installing it via a docker container was always the goal. So lets proceed with option 1.

For the things you asked:

- platform.ini:

```
; PlatformIO Project Configuration File
;
; Build options: build flags, source filter
; Upload options: custom upload port, speed and extra flags
; Library options: dependencies, extra library storages
; Advanced options: extra scripting
;
; Please visit documentation for the other options and examples
; https://docs.platformio.org/page/projectconf.html

[env:teensy41]
platform = teensy
board = teensy41
framework = arduino
monitor_speed = 115200
board_microros_distro = humble
board_microros_transport = serial
lib_deps = https://github.com/micro-ROS/micro_ros_platformio
```

- error when using humble:

```
Downloading micro-ROS dev dependencies
ament_cmake pull failed:
From https://github.com/ament/ament_cmake
* branch          humble      -> FETCH_HEAD
hint: You have divergent branches and need to specify how to reconcile them.
hint: You can do so by running one of the following commands sometime before
hint: your next pull:
hint:
hint:   git config pull.rebase false  # merge
hint:   git config pull.rebase true   # rebase
hint:   git config pull.ff only       # fast-forward only
hint:
hint: You can replace "git config" with "git config --global" to set a default
hint: preference for all repositories. You can also pass --rebase, --no-rebase,
hint: or --ff-only on the command line to override the configured default per
hint: invocation. fatal: Need to specify how to reconcile divergent branches.
```

- System:

debian 12 (at the moment), but this needs to run on an nvidia jetson nano afterwards (this can be ignored for now)

Note that i'm storing this all in my own git repo: TruggySense And that the platformio project should be located on
 "~/Documents/studies/4_UA/semester_8/II-masterproef_elektronica_ict_met_inbegrip_van_stage/TruggySense/code/protocol_tests/ROS2"

If possible could you link this project to the one inside the docker container, that is going to program the teensy?

The teensy is connected via the /dev/ttyACM0 port (for reference)

LLM response

Perfect — now we're getting to the real issue 🙌

You've done things mostly right, but you hit a **known micro-ROS + PlatformIO + git bug** inside Docker.

Let's fix this cleanly and set up a **proper Docker-based workflow that builds + flashes your Teensy**.

Problem you're hitting now

This error:

fatal: Need to specify how to reconcile divergent branches.

➞ Comes from **git inside the micro-ROS build system**, not your code.

Why it happens

- micro-ROS tries to `git pull` dependencies (like `ament_cmake`)
- Newer Git (Debian 12) **requires explicit pull strategy**
- The script doesn't set it → build fails

✓ Fix (simple and reliable)

Inside your Docker environment, run:

```
git config --global pull.rebase false
```

➞ This fixes the error permanently.

🦋 Now let's build your FULL workflow (clean & working)

We'll do exactly what you want:

- ✓ Your project on host
- ✓ Docker builds + uploads firmware
- ✓ Uses `/dev/ttyACM0`
- ✓ No local ROS install needed

Step 1 — Create a proper Dockerfile for PlatformIO + micro-ROS

Create this in your repo (e.g. `Dockerfile.pio`):

```
FROM ros:humble
```

```
SHELL ["/bin/bash", "-c"]
```

```
# Install dependencies
```

```
RUN apt update && apt install -y \
    python3-pip \
    git \
    build-essential \
    && rm -rf /var/lib/apt/lists/*
```

```
# Fix git issue (IMPORTANT)
RUN git config --global pull.rebase false
```

```
# Install PlatformIO
RUN pip install platformio
```

```
WORKDIR /workspace
```

```
CMD ["/bin/bash"]
```

Step 2 — Build the container

From your project root:

```
docker build -t microros-pio -f Dockerfile.pio .
```

Step 3 — Run container with your project mounted

```
docker run -it --rm \
  --privileged \
  -v ~/Documents/studies/4_UA/semester_8/II-masterproef_elektronica_ict_met_inbegrip_van_stage/TruggySense/code/protocol_tests/ROS2:/workspace \
  -v /dev:/dev \
  microros-pio
```

⇒ Now inside container:

- Your PlatformIO project is at /workspace
- Teensy is accessible via /dev/ttyACM0

Step 4 — Build & upload (inside container)

```
cd /workspace
pio run -t upload
```

What just happened

Now:

- micro-ROS builds inside **ROS 2 Humble environment** ✓
- rmw exists ✓
- git bug fixed ✓
- Teensy gets flashed ✓

One more fix you MAY need

If Teensy upload fails, install udev rules:

```
apt update
apt install -y udev
```

Sometimes also needed:

```
chmod 666 /dev/ttyACM0
```

Step 5 — Run the agent (separately)

On your host:

```
docker run -it --rm \
  --privileged \
  --net=host \
  -v /dev:/dev \
  microros/micro-ros-agent:humble \
  serial --dev /dev/ttyACM0 -v6
```

Architecture you now have

```
[ Teensy 4.1 ]
  | USB (/dev/ttyACM0)
  ▼
[ micro-ROS agent (Docker) ]
  ▼
[ ROS 2 DDS network ]
```

And separately:

```
[ PlatformIO build container ]
  └─ compiles + uploads firmware
```

💡 Optional (nice improvement)

docker-compose version

You can automate both containers:

version: "3.8"

```
services:
  builder:
    build:
      context: .
      dockerfile: Dockerfile.pio
    volumes:
      - ./code/protocol_tests/ROS2:/workspace
      - /dev:/dev
    privileged: true
    stdin_open: true
    tty: true

  agent:
    image: microros/micro-ros-agent:humble
    command: serial --dev /dev/ttyACM0 -v6
    network_mode: host
    privileged: true
    volumes:
      - /dev:/dev
```



Summary (what fixed everything)

- ✓ Build inside ROS-enabled container
- ✓ Add git config --global pull.rebase false

- ✓ Mount your project into container
- ✓ Use official agent separately

If something still breaks

Tell me:

- Output of `pio run` (after this fix)
- Whether upload fails or build fails

I can debug the exact micro-ROS stage (it has a few 😊)

LLM request

When doing step 4 (running the "pio run -t upload" command in the container), i get this error:

```
Configuring teensy41 with transport serial
Downloading micro-ROS dev dependencies
ament_cmake pull failed:
fatal: detected dubious ownership in repository at
'/workspace/.pio/libdeps/teensy41/micro_ros_platformio/build/dev/src/ament_cmake'
To add an exception for this directory, call:

git config --global --add safe.directory
/workspace/.pio/libdeps/teensy41/micro_ros_platformio/build/dev/src/ament_cmake
```

Is there a solution for this?

LLM response

Yep — this is another **modern Git safety feature** biting Docker workflows 😊

Good news: it's easy to fix once you understand it.

What's happening

Git is complaining:

fatal: detected dubious ownership in repository

⇒ कारण:

- Your project is mounted from host → owned by your host user (UID ≠ root)
- Inside Docker, you're running as **root**
- Git sees a mismatch → blocks access for security

This is expected on **Debian 12 / newer Git versions**

✓ Quick fix (works immediately)

Inside the container, run:

```
git config --global --add safe.directory '*'
```

⇒ This tells Git:

"I trust all mounted directories"

Then rerun:

```
pio run -t upload
```

✓ Better fix (persist in Dockerfile)

Update your Dockerfile.pio:

```
# Fix git issues
```

```
RUN git config --global pull.rebase false && \
    git config --global --add safe.directory '*'
```

Then rebuild:

```
docker build -t microros-pio -f Dockerfile.pio .
```

🧠 Why this keeps happening in micro-ROS

micro-ROS does:

- clones repos into .pio/libdeps/...
- runs `git pull` inside them

So Git checks ownership multiple times → fails unless trusted

🔒 Alternative (more secure, optional)

Instead of trusting everything:

```
git config --global --add safe.directory /workspace
```

But micro-ROS creates nested repos, so you'd need multiple entries → annoying.

⇒ That's why '*' is commonly used in containers.



After this fix

Your flow should now:

1. ✓ Pull dependencies
2. ✓ Build micro-ROS stack
3. ✓ Compile firmware
4. ✓ Upload to Teensy



If you hit the NEXT error (likely)

If something else breaks after this, it's usually:

✗ Missing Python deps

Fix:

```
pip install empy colcon-common-extensions
```

✗ Teensy upload fails

Fix:

```
chmod 666 /dev/ttyACM0
```



Summary

Add this once:

```
git config --global --add safe.directory '*'
```

or bake it into Dockerfile → and you're past the last major blocker.

If the build proceeds but fails later, paste the next error — you're now in the “final 10%” of setup 👉

LLM request

platformio doesn't seem to create

```
~/platformio/penv
```

when running

```
pio system info
```

LLM response

Ah! I see what's happening. PlatformIO should normally create a Python virtual environment under `~/.platformio/penv` when it needs to run certain commands, but there are a few subtleties that can prevent it from appearing. Let's go step by step.

1 Why `~/.platformio/penv` might not exist

`pio system info` itself **doesn't always trigger the creation of the virtual environment**. The `penv` folder is only created when PlatformIO needs its **isolated Python environment** to install packages or dependencies. For example:

- Installing a platform or library.
- Running `pio upgrade` (sometimes).
- Using `pio run` with certain packages.

So just checking `pio system info` may **not yet require the penv**, especially if your PlatformIO installation is already satisfied with your system Python.

2 Check your installation type

PlatformIO can be installed in a few ways:

- **Homebrew / apt / pacman system package**
- **Python pip package**
- **VSCode extension (PIO IDE)**

Only the **pip-based installation** usually manages `~/.platformio/penv`. If you installed via a system package manager, PlatformIO may rely on your system Python instead and skip the `penv`.

You can check your installation path with:

```
which pio
pio --version
```

3 Force creation of `penv`

If you want `~/.platformio/penv` to exist, you can trigger it by asking PlatformIO to **reinstall its virtual environment**:

```
pio system install --force
```

Or, if you want to manually ensure a virtual environment is created:

```
python3 -m venv ~/.platformio/penv
~/.platformio/penv/bin/python -m pip install -U platformio
```

This is usually **not necessary** unless you need to debug Python dependency issues.

4 Confirm virtual environment usage

Once `penv` exists, you can activate it:

```
source ~/.platformio/penv/bin/activate
pio --version
```

This ensures `pio` uses the isolated environment.

✓ Summary:

- `~/.platformio/penv` is only created when PlatformIO needs its own isolated Python environment.
- `pio system info` alone might not trigger it.
- Use `pio system install --force` to create it if needed.
- Check your installation method—system package vs pip—because it changes

where penv appears.

If you want, I can give you a **single command to fully recreate penv and ensure PlatformIO uses it**, which usually fixes any “missing virtual environment” issues.

Do you want me to do that?

LLM request

I still get these errors:

```

root@8ff623b2e945:/workspace# rm -rf /root/.platformio
pio system info
*****
*****
*****
If you like PlatformIO, please:
- star it on GitHub > https://github.com/platformio/platformio-core
- follow us on LinkedIn to stay up-to-date on the latest project news >
https://www.linkedin.com/company/platformio/
- try PlatformIO IDE for embedded development > https://platformio.org/platformio-
ide
*****
*****
*****

-----
PlatformIO Core 6.1.19
Python 3.10.12-final.0
System Type linux_x86_64
Platform Linux-6.1.0-43-amd64-x86_64-with-glibc2.35
File System Encoding utf-8
Locale Encoding UTF-8
PlatformIO Core Directory /root/.platformio
PlatformIO Core Executable /usr/local/bin/platformio
Python Executable /usr/bin/python3
Global Libraries 0
Development Platforms 0
Tools & Toolchains 0
-----

root@8ff623b2e945:/workspace# ls /root/.platformio appstate.json

root@8ff623b2e945:/workspace# pio system install --force
Usage: pio system [OPTIONS] COMMAND [ARGS]...
Try 'pio system -h' for help.
Error: No such command 'install'.

root@8ff623b2e945:/workspace# ls /root/.platformio appstate.json

root@8ff623b2e945:/workspace# which pio
pio --version
/usr/local/bin/pio
PlatformIO Core, version 6.1.19

```

```
root@8ff623b2e945:/workspace# pio system install --force
Usage: pio system [OPTIONS] COMMAND [ARGS]...
Try 'pio system -h' for help.
Error: No such command 'install'.

root@8ff623b2e945:/workspace# pio lib -g install --force
WARNING: This command is deprecated and will be removed in the next releases.
Please use pio pkg install instead.

root@8ff623b2e945:/workspace# pio pkg -g install --force
Usage: pio pkg [OPTIONS] COMMAND [ARGS]...
Try 'pio pkg -h' for help. Error: No such option: -g

root@8ff623b2e945:/workspace# pio pkg install --force
Resolving teensy41 dependencies...
Platform Manager: Installing teensy
Downloading [#####] 100%
Unpacking [#####] 100%
Platform Manager: teensy@5.1.0 has been installed!
Tool Manager: Installing platformio/toolchain-gccarmnoneabi-teensy @ ~1.110301.0
Downloading [#####] 100%
Unpacking [#####] 100%
Tool Manager: toolchain-gccarmnoneabi-teensy@1.110301.0 has been installed!
Tool Manager: Installing platformio/framework-arduinoteensy @ ~1.160.0
Downloading [#####] 100%
Unpacking [#####] 100%
Tool Manager: framework-arduinoteensy@1.160.0 has been installed!
Tool Manager: Installing platformio/tool-teensy @ <2
Downloading [#####] 100%
Unpacking [#####] 100%
Tool Manager: tool-teensy@1.159.0 has been installed!
Tool Manager: Installing platformio/tool-scons @ ~4.40801.0
Downloading [#####] 100%
Unpacking [#####] 100%
Tool Manager: tool-scons@4.40801.0 has been installed!
Library Manager: Removing micro_ros_platformio @ 0.0.1+sha.cfee17f
Library Manager: micro_ros_platformio@0.0.1+sha.cfee17f has been removed!
Library Manager: Installing git+https://github.com/micro-ROS/micro_ros_platformio
git version 2.34.1
Cloning into '/root/.platformio/.cache/tmp/pkg-installing-fgu81g__'...
remote: Enumerating objects: 77, done.
remote: Counting objects: 100% (77/77), done.
remote: Compressing objects: 100% (62/62), done.
remote: Total 77 (delta 12), reused 38 (delta 4), pack-reused 0 (from 0)
Receiving objects: 100% (77/77), 132.22 KiB | 2.75 MiB/s, done.
Resolving deltas: 100% (12/12), done.
Library Manager: micro_ros_platformio@0.0.1+sha.cfee17f has been installed!
```



```
root@8ff623b2e945:/workspace# ls
/root/.platformio appstate.json packages platforms
```

```
root@8ff623b2e945:/workspace# pio platform install --force
```

```
Usage: pio platform install [OPTIONS] [PLATFORM...]
```

```
Try 'pio platform install -h' for help.
```

```
Error: Missing argument '[PLATFORM...]'.

```

```
root@8ff623b2e945:/workspace# python3 -m venv ~/.platformio/penv
```

```
~/.platformio/penv/bin/python -m pip install -U platformio
```

```
The virtual environment was not created successfully because ensurepip is not
available. On Debian/Ubuntu systems, you need to install the python3-venv package
using the following command.
```

```
apt install python3.10-venv
```

You may need to use sudo with that command. After installing the python3-venv package, recreate your virtual environment.

Failing command: /root/.platformio/penv/bin/python3

```
/root/.platformio/penv/bin/python: No module named pip
```

```
root@8ff623b2e945:/workspace# python3 -m venv ~/.platformio/penv
```

```
The virtual environment was not created successfully because ensurepip is not
available. On Debian/Ubuntu systems, you need to install the python3-venv package
using the following command.
```

```
apt install python3.10-venv
```

You may need to use sudo with that command. After installing the python3-venv package, recreate your virtual environment.

Failing command: /root/.platformio/penv/bin/python3

```
root@8ff623b2e945:/workspace# apt install python3.10-venv
```

```
Reading package lists...
```

```
Done Building dependency tree...
```

```
Done Reading state information...
```

```
Done Package python3.10-venv is not available, but is referred to by another package.
This may mean that the package is missing, has been obsoleted, or is only available
from another source
```

```
E: Package 'python3.10-venv' has no installation candidate
```

```
root@8ff623b2e945:/workspace#
```

LLM response

Ah! Now everything makes sense — the core issue is **your environment doesn't have the Python venv module installed**, so PlatformIO can't create `~/.platformio/venv`. Let's break it down carefully and fix it.

1 Why venv isn't created

- You're running **Debian/Ubuntu 6.1 container / Linux x86_64**.
- Your Python 3.10 installation lacks the venv module (ensurepip), which is required to create virtual environments.
- Commands like:

```
python3 -m venv ~/.platformio/venv
```

fail because python3-venv isn't installed.

2 Fix missing venv on Debian/Ubuntu 3.x / containers

You tried:

```
apt install python3.10-venv
```

...but it failed because your package source probably **doesn't have a package for the exact Python 3.10 build** in your container. That's very common in minimal Docker images.

Instead, do this:

```
apt update
```

```
apt install python3-venv python3-pip -y
```

Explanation:

- python3-venv is the generic package for the default Python 3 version (usually matches python3.10).
- python3-pip ensures pip is available inside virtual environments.

3 Recreate the PlatformIO virtual environment

After installing python3-venv:

```
python3 -m venv ~/.platformio/venv
```

```
~/.platformio/venv/bin/python -m pip install --upgrade pip
```

```
~/.platformio/venv/bin/python -m pip install -U platformio
```

✓ This will create `~/.platformio/venv` properly.

4 Confirm everything works

```
ls ~/.platformio/venv
```

```
source ~/.platformio/venv/bin/activate
```

```
pio --version
```

You should see PlatformIO Core 6.1.19 running inside its isolated Python environment.

LLM request

When i run this, i get the following error:

```
Building in release mode
Compiling .pio/build/teensy41/src/main.cpp.o
Compiling .pio/build/teensy41/libc8d/micro_ros_platformio/platform_code/arduino/clock_gettime.cpp.o
Compiling .pio/build/teensy41/libc8d/micro_ros_platformio/platform_code/arduino/serial/micro_ros_transport.cpp.o
Compiling .pio/build/teensy41/FrameworkArduino/AudioStream.cpp.o
Compiling .pio/build/teensy41/FrameworkArduino/CrashReport.cpp.o
Compiling .pio/build/teensy41/FrameworkArduino/DMAChannel.cpp.o
src/main.cpp: In function 'void setup()': src/main.cpp:14:3: error: 'Serial' was not declared in this scope
14 |   Serial.begin(115200);
   |   ^~~~~~
src/main.cpp:15:3: error: 'delay' was not declared in this scope
15 |   delay(2000);
   |   ^~~~~
src/main.cpp: In function 'void loop()':
src/main.cpp:37:3: error: 'delay' was not declared in this scope
37 |   delay(1000);
   |   ^~~~~
src/main.cpp:36:14: warning: ignoring return value of 'rcl_ret_t rcl_publish(const rcl_publisher_t*, const void*, rmw_publisher_allocation_t*)' declared with attribute 'warn_unused_result' [-Wunused-result]
36 |   rcl_publish(&publisher, &msg, NULL);
   |   ~~~~~~^~~~~~
*** [.pio/build/teensy41/src/main.cpp.o]
Error 1 Compiling .pio/build/teensy41/FrameworkArduino/EventResponder.cpp.o
```

LLM response

Thanks for providing all the details — this helps a lot. From the error you're getting:
src/main.cpp:14:3: error: 'Serial' was not declared in this scope
src/main.cpp:15:3: error: 'delay' was not declared in this scope
This is a classic **Arduino/PlatformIO issue**, and it usually comes down to **missing #include <Arduino.h> in your main.cpp**. On Teensy boards, Serial, delay(), millis(), etc., are all defined in Arduino.h, and unlike the Arduino IDE, PlatformIO **does not automatically include it** for you.

So the fix is simple: at the top of src/main.cpp, add:

```
#include <Arduino.h>
```

For example:

```
#include <Arduino.h>
#include <rcl/rcl.h>
#include <std_msgs/msg/int32.h>
// any other includes
```

```
void setup() {
    Serial.begin(115200);
    delay(2000);
}
```

```
void loop() {
    // your loop code
    delay(1000);
}
```

LLM request

I get this error from inside the container

```
Uploading .pio/build/teensy41/firmware.hex  
Opening Teensy Loader...  
Unable find Teensy Loader. (p) Is the Teensy Loader application running?  
quitscons: *** [upload] Error 1
```

LLM response

Ah — now we're hitting a **Teensy upload issue inside Docker**, and it's not a code problem. Let me break this down carefully:

What's happening

PlatformIO uses **Teensy Loader CLI** (`teensy_loader_cli`) to upload firmware. The error:

Unable find Teensy Loader. (p) Is the Teensy Loader application running?

means **PlatformIO cannot access the Teensy Loader**, either because:

1. The CLI is not installed inside your Docker container.
2. The container cannot access the host GUI Teensy Loader app.
3. The container cannot see the `/dev/ttyACM0` (Teensy USB port).

Why this happens in Docker

- The Teensy Loader GUI is normally installed on your **host system** (Debian 12).
- Docker containers are isolated — they **cannot access GUI apps** by default.
- `--privileged -v /dev:/dev` gives raw USB access, but PlatformIO **still expects the Teensy Loader CLI**.

How to fix

Option 1 — Install `teensy_loader_cli` inside the container (recommended)

1. Inside your Dockerfile, add:

```
# Install Teensy Loader CLI
RUN apt update && apt install -y wget unzip \
    && wget \
    https://www.pjrc.com/teensy/td_266/TeensyduinoCLI_2.6.6_Linux.tar.gz -O /tmp/teensycli.tar.gz \
    && tar -xzf /tmp/teensycli.tar.gz -C /usr/local/bin --strip-components=1 \
    && rm /tmp/teensycli.tar.gz
```

Note: Replace the URL with the latest Teensy Loader CLI for Linux if needed.

1. Make sure your container has `teensy_loader_cli` in `$PATH`:

```
which teensy_loader_cli
```

It should print something like `/usr/local/bin/teensy_loader_cli`.

Option 2 — Use host Teensy Loader

- Upload firmware **from your host** (Debian 12), not inside Docker.
- In that case, do **pio run to build**, then **use Teensy Loader GUI or CLI** on the host to flash `.pio/build/teensy41/firmware.hex`.

Option 3 — Set `upload_protocol = teensy-cli` explicitly

Add to `platformio.ini`:

```
upload_protocol = teensy-cli
```

This tells PlatformIO to **look for the CLI**, not the GUI. Combined with Option 1, it will work fully inside Docker.

TL;DR Recommendation

For a fully Dockerized Teensy + PlatformIO + micro-ROS workflow:

1. Install `teensy_loader_cli` inside your container.
2. Set `upload_protocol = teensy-cli` in `platformio.ini`.

3. Run the container with `--privileged -v /dev:/dev` so USB access works.

4. Use:
`pio run -t upload -e teensy41`
inside the container. It should upload successfully.

I can give you a **ready-to-use Dockerfile snippet** that adds Teensy Loader CLI and PlatformIO fully for Teensy 4.1 + micro-ROS, so your upload will work without touching the host.

LLM request

Option 3 worked to upload the the sketch, but the serial output gave this:

```
~\...\XRCE^\L?h
```


LLM response

Ah! That “garbled” output is actually **very typical for micro-ROS over Serial** — it’s not a code error. Let me explain carefully.

Why it looks like garbage

- micro-ROS on **Teensy** uses the **XRCE-DDS protocol**, which is a binary protocol for ROS 2 communication.
- When you `Serial.begin(115200)` and then open the serial monitor, you're seeing **raw binary XRCE frames**, not human-readable text.
- That explains why you get output like:

$$\sim N_{U_L} N_{U_L} C_A N_{U_L} \diamond N_{U_L} N_{U_L} N_{U_L} N_{U_L} S_O H^D E N_{U_L} X R C E S_O H^S N_{U_L} S_O H^S I \} \wedge L \diamond \diamond N_{U_L} \diamond S_O H \diamond h$$

This is expected. The Teensy is **sending ROS messages**, not ASCII logs.

How to correctly read micro-ROS data

You **don't** use a **regular serial monitor** for these messages. You have to use a **micro-ROS agent** running on your host machine that connects to the Teensy over serial.

For example:

1. Install micro-ROS agent on your host (Debian 12):

```
sudo apt update
```

```
sudo apt install python3-pip
```

```
pip3 install -U micro_ros_agent
```

1. Run the agent for a serial transport:

```
ros2 run micro_ros_agent micro_ros_agent serial --dev /dev/ttyACM0 -b 115200
```

- /dev/ttyACM0 → your Teensy serial port
- -b 115200 → must match Serial.begin(115200) on Teensy

1. Then, on your host, you can **inspect published messages** with standard ROS 2 tools:

```
ros2 topic echo /teensy_counter
```

You should see messages like:

```
data: 1
```

— — —

```
data: 2
```

— — —

```
data: 3
```

— — —

Optional: Debug logs from Teensy

If you want to **see debug prints on the Teensy** while still using micro-ROS, you can:

```
// Only for debugging
```

```
SerialUSB.println("Hello micro-ROS");
```

...but remember **micro-ROS over serial and debug prints share the same port**, so mixing them can **break the XRCE protocol**. A safer way is to:

- Use **separate USB ports** if your board supports it
- Or send debug info via **LEDs or another channel**

✓ Summary

- The “garbled” serial output is **normal XRCE-DDS frames**.
- You must run **micro-ROS Agent** on the host to interpret it.

- Use `ros2 topic echo` to see the published messages.

If you want, I can **rewrite your Teensy sketch** so it **blinks the LED and publishes micro-ROS messages at exactly 1 Hz**, while keeping the serial free for micro-ROS only. This is the cleanest setup for testing.

Do you want me to do that?

LLM request

I want to keep things inside the container

LLM response

Perfect — we can keep **everything inside Docker**, including building, uploading, and running the micro-ROS agent. Here's how you can do it correctly.

1 Update Dockerfile to include micro-ROS Agent

Add this to your Dockerfile after installing Python:

```
# Install micro-ROS Agent
RUN pip install -U micro-ros-agent
```

Now your container can **run the agent directly**.

2 Make sure Teensy serial is accessible

You already run the container with:

```
docker run -it --rm --privileged -v /dev:/dev ...
```

This gives the container access to /dev/ttyACM0 (Teensy).

Inside the container, you can verify:

```
ls /dev/ttyACM*
```

It should show your Teensy port.

3 Start the micro-ROS Agent inside the container

Instead of opening the serial monitor, run:

```
ros2 run micro_ros_agent micro_ros_agent serial --dev /dev/ttyACM0 -b 115200
```

- /dev/ttyACM0 → Teensy port inside container
- -b 115200 → matches Serial.begin(115200) in your sketch

You should see the agent start:

```
[INFO] --- micro-ROS Agent initialized ---
```

Now your Teensy is connected to the ROS 2 network inside the container.

4 Inspect messages

To see the messages, you can run another terminal inside the same container:

```
docker exec -it <container_name_or_id> /bin/bash
```

Then, inside the container:

```
ros2 topic echo /teensy_counter
```

You should see:

```
data: 1
```

```
---
```

```
data: 2
```

```
---
```

5 Optional: LED + message publish at 1 Hz

If you want the LED blink and ROS message independent of each other:

```
#include <Arduino.h>
#include <micro_ros_platformio.h>
#include <rcl/rcl.h>
#include <rcl/rcl.h>
#include <rcl/rcl.h>
#include <std_msgs/msg/int32.h>
```

```
#ifndef LED_BUILTIN
#define LED_BUILTIN 13
#endif
```

```
rcl_publisher_t publisher;
```

```

rcl_node_t node;
rcl_support_t support;
rcl_allocator_t allocator;
std_msgs__msg__Int32 msg;

unsigned long previousMillis = 0;
const long interval = 1000; // 1 second

void setup() {
  Serial.begin(115200);
  delay(2000);
  pinMode(LED_BUILTIN, OUTPUT);

  set_microros_serial_transports(Serial);

  allocator = rcl_get_default_allocator();
  rcl_support_init(&support, 0, NULL, &allocator);
  rcl_node_init_default(&node, "teensy_node", "", &support);
  rcl_publisher_init_default(
    &publisher, &node,
    ROSIDL_GET_MSG_TYPE_SUPPORT(std_msgs, msg, Int32),
    "teensy_counter"
  );

  msg.data = 0;
}

void loop() {
  unsigned long currentMillis = millis();

  // Blink LED every 500 ms
  if (currentMillis % 1000 < 500) {
    digitalWrite(LED_BUILTIN, HIGH);
  } else {
    digitalWrite(LED_BUILTIN, LOW);
  }

  // Publish message every 1 second
  if (currentMillis - previousMillis >= interval) {
    previousMillis = currentMillis;
    msg.data++;
    rcl_publish(&publisher, &msg, NULL);
  }
}

```

- LED blinks at 1 Hz
- micro-ROS messages publish at 1 Hz **independently**

LLM request

```
pip install -U micro-ros-agent
ERROR: Could not find a version that satisfies the requirement micro-ros-agent (from
versions: none)
ERROR: No matching distribution found for micro-ros-agent
```

LLM response

Right — the reason you got this error when trying to install `micro_ros_agent` with `pip` is simply because **there is no Python package on PyPI called `micro-ros-agent`**. It's *not* a `pip`-installable package.

The **micro-ROS Agent** is a ROS-2 node built from source inside a ROS-2 workspace, and *not* a Python package you can install via `pip install ...`

✓ What the micro-ROS Agent *actually* is

The micro-ROS Agent is a **ROS-2 executable** that bridges your micro-ROS device (over serial, UDP, etc.) to the ROS-2 graph on your host (or container). You normally build it either by:

● Option A — Using `micro_ros_setup` toolchain

This is the official supported way inside a ROS-2 environment:

```
# inside a ROS2 workspace
ros2 run micro_ros_setup create_agent_ws.sh
ros2 run micro_ros_setup build_agent.sh
source install/local_setup.bash
ros2 run micro_ros_agent micro_ros_agent serial --dev /dev/ttyACM0 -b 115200
```

This fetches the Agent source and builds it as a regular ROS-2 package.

● Option B — Use Official Docker Image

Instead of building the Agent yourself, you can run an **official micro-ROS Agent Docker image**, which already contains the compiled agent binary. For example:

```
docker run -it --rm \
  --privileged \
  -v /dev:/dev \
  --net=host \
  microros/micro-ros-agent:humble
```

Then inside the container:

```
micro_ros_agent serial --dev /dev/ttyACM0 -b 115200
```

This is often the simplest approach because the agent is already pre-built.

⚠ Important: micro-ROS Agent $\neq$ Python `pip` package

- There **is no `micro_ros_agent` on PyPI** to install with `pip`.
- The Agent is a **ROS-2 node**, not a Python lib.
- The usual way to install it is via ROS-2 build tools or via a Docker image.

🔑 Running the Agent inside *your* Dev Container

Since you want everything inside your **PlatformIO / ROS container**, you have three realistic options:

✓ Option 1 — Build agent inside your container (requires ROS-2 tooling)

Inside your container (which is based on `ros:humble`), install:

```
sudo apt update
sudo apt install python3-rosdep build-essential
rosdep update
```

Then use `micro_ros_setup` to create and build the agent:

```
mkdir -p ~/agent_ws/src
cd ~/agent_ws
ros2 run micro_ros_setup create_agent_ws.sh
rosdep install --from-paths src --ignore-src -y
colcon build
source install/local_setup.bash
```

Then run:

```
ros2 run micro_ros_agent micro_ros_agent serial --dev /dev/ttyACM0 -b 115200
```

This keeps everything in the same container as your build environment.

✓ Option 2 — Use the existing official agent Docker image

You can run a *separate container* just for the agent:

```
docker run -it --rm \
  --privileged \
  -v /dev:/dev \
  --net=host \
  microros/micro-ros-agent:humble \
  serial --dev /dev/ttyACM0 -b 115200
```

Then in another container you can do `ros2 topic echo /teensy_counter`.

🔑 Notes:

- **You still need a ROS-2 installation** inside the container to run `ros2 topic echo`, `ros2 node list`, etc.
- The agent binary is part of ROS-2 workspace builds — not pip.
- Snap packages *do* exist for micro-ROS Agent on some systems, but not via pip, and not inside Docker without snapd.

✂ Summary

Task	How
Install micro-ROS Agent <code>pip install micro-ros-agent</code>	Build via <code>micro_ros_setup</code> OR use official Docker image ✗ Not possible — no package on PyPI
Run agent inside container	✓ Use ROS-2 tools or official <code>microros/micro-ros-agent</code> image
Connect Teensy to ROS	✓ <code>micro_ros_agent serial --dev ... -b 115200</code>

If you want, I can give you a **ready-to-paste Dockerfile** that already includes ROS-2, PlatformIO, Teensy CLI, and the micro-ROS Agent built and installed — so you can do *everything* in one container. Just ask!

LLM request

I'm inside that official docker image, remember.

```
FROM ros:humble
```

LLM response

Ah — thanks for clarifying! Yes, your **Dockerfile is FROM ros:humble**, but that is **the standard ROS 2 Humble image**, not the **official micro-ROS Agent image**.

That's why `pip install micro-ros-agent` fails — the standard `ros:humble` image **does not include the micro-ROS Agent binary**, and it isn't a Python package anyway.

So right now your container can:

- Build and upload your Teensy micro-ROS firmware (via PlatformIO + Teensy CLI). ✓
- But it **cannot run `micro_ros_agent`** yet. ✗

How to run micro-ROS Agent inside your current container

You have two options:

Option 1 — Build micro-ROS Agent inside your container

Since your container already has ROS 2 Humble:

1. Create a workspace

```
mkdir -p ~/microros_agent_ws/src
```

```
cd ~/microros_agent_ws
```

2. Clone `micro_ros_setup`

```
git clone -b humble https://github.com/micro-ROS/micro-ros-build.git
src/micro_ros_setup
```

3. Install `colcon` if not present

```
apt update && apt install -y python3-colcon-common-extensions
```

4. Install dependencies

```
rosdep update
```

```
rosdep install --from-paths src --ignore-src -y
```

5. Build the workspace

```
colcon build
```

6. Source workspace

```
source install/local_setup.bash
```

7. Run the agent (replace `ttyACM0` with your Teensy)

```
ros2 run micro_ros_agent micro_ros_agent serial --dev /dev/ttyACM0 -b 115200
```

✓ This will give you the `micro_ros_agent` executable in your container.

Option 2 — Use the prebuilt official micro-ROS Agent Docker image

Instead of building inside your own container, you could simply run:

```
docker run -it --rm \
  --privileged \
  -v /dev:/dev \
  --net=host \
  microros/micro-ros-agent:humble \
  serial --dev /dev/ttyACM0 -b 115200
```

- This container already has **ROS 2 Humble + micro-ROS Agent binary**.
- You only need your Teensy plugged into `/dev/ttyACM0`.

- Then you can either run `ros2 topic echo /teensy_counter` inside the same container (or another ROS 2 container).

✓ Recommendation

Since you already have a **PlatformIO + ROS 2 container**, the **fastest way** is **Option 1**, build the micro-ROS Agent inside your container. That keeps **everything inside one Docker container**, which is exactly what you wanted.

If you want, I can give you a **full Dockerfile example** that:

1. Starts from `ros:humble`
2. Installs PlatformIO & Teensy CLI
3. Builds your micro-ROS Agent
4. Is ready to build, upload, and run micro-ROS + Agent entirely inside the container

Do you want me to do that?

LLM request

```
ros2 run micro_ros_setup create_agent_ws.sh  
Package 'micro_ros_setup' not found
```

LLM response

Ah — now we're hitting a very common **confusion with micro-ROS tooling**. Let me clarify carefully.

Why `ros2 run micro_ros_setup create_agent_ws.sh` fails

- `micro_ros_setup` is **not** a ROS 2 package installed by default.
- It's actually a **repository you have to clone into a workspace** before you can run its scripts.
- That's why ROS 2 says: Package 'micro_ros_setup' not found.

Correct way to set up micro-ROS Agent inside your container

1 Clone `micro_ros_setup`

Inside your container:

```
# Create workspace
mkdir -p ~/microros_agent_ws/src
cd ~/microros_agent_ws/src

# Clone micro_ros_setup for Humble
git clone -b humble https://github.com/micro-ROS/micro-ros-build.git
micro_ros_setup
```

2 Install ROS dependencies

```
# Go to workspace root
cd ~/microros_agent_ws

# Update rosdep
rosdep update

# Install missing dependencies
rosdep install --from-paths src --ignore-src -y
```

3 Build the workspace with `colcon`

```
# Make sure colcon is installed
apt update && apt install -y python3-colcon-common-extensions

# Build
colcon build
```

4 Source the workspace

```
source install/local_setup.bash
Now micro_ros_setup scripts are available as executable scripts.
```

5 Create and build the Agent

```
# Go to workspace root again
cd ~/microros_agent_ws

# Run agent setup script
ros2 run micro_ros_setup create_agent_ws.sh

# Install any missing dependencies (if prompted)
rosdep install --from-paths src --ignore-src -y

# Build the Agent
ros2 run micro_ros_setup build_agent.sh
```

```
# Source the agent
source install/local_setup.bash
After this, you can run:
ros2 run micro_ros_agent micro_ros_agent serial --dev /dev/ttyACM0 -b
115200
```

✓ Key Points

1. micro_ros_setup is **not pre-installed**, you must clone it.
2. You **cannot pip install micro-ros-agent**. It's a ROS 2 node built from source.
3. After building inside the container, the micro_ros_agent executable is ready, so you can **stay entirely inside Docker**.
4. Make sure /dev/ttyACM0 is accessible to the container (--privileged -v /dev:/dev).

I can make you a **ready-to-use Dockerfile** that:

- Starts from ros:humble
- Installs PlatformIO + Teensy CLI
- Clones and builds micro_ros_setup and the micro-ROS Agent

...so you can immediately build, upload, and run micro-ROS **all inside one container**.

Do you want me to do that?

Something is still going wrong

Page 55 of 70

LLM response

Ah! Perfect — now this tells the full story.

What you are seeing is **exactly how micro-ROS is supposed to work**. Let me break it down:

What's happening

1. **Teensy is publishing XRCE-DDS frames**
2. When you ran:
3. `pio device monitor -b 115200`
4. you saw:
5. `~^LXRCE}^Lh`
6. This is **binary micro-ROS data**, not meant to be human-readable. This is completely normal.
7. **micro-ROS Agent running**
8. You then ran:
9. `ros2 run micro_ros_agent micro_ros_agent serial --dev /dev/ttyACM0 -b 115200`
10. And you see:
11. `[1774011836.240819] info | TermiosAgentLinux.cpp | init | running... | fd: 3`
12. ✓ This means the agent is successfully connected to your Teensy over serial. The agent is now bridging the Teensy to the ROS 2 network **inside the container**.

What to do next

Now, to actually see your messages being published by the Teensy:

1. **Open a new shell inside the container** (don't stop the agent):
2. `docker exec -it <container_id_or_name> /bin/bash`
`source /opt/ros/humble/setup.bash` # Make sure ROS2 environment is sourced
3. **Echo the topic published by Teensy:**
4. `ros2 topic list`
5. You should see something like:
6. `/teensy_counter`
7. Then:
8. `ros2 topic echo /teensy_counter`
9. Output should look like:
10. `data: 1`
`---`
`data: 2`
`---`
`data: 3`
`---`
11. That's your Teensy publishing micro-ROS messages via the agent.

✓ Key points

- The “garbled serial output” is normal — don't use `pio device monitor` to

read micro-ROS messages; it will always be binary.

- **micro-ROS Agent translates the binary messages** into standard ROS 2 topics.
- You now have a **fully working setup inside Docker**: build → upload → run agent → subscribe/publish.

If you want, I can provide a **step-by-step Docker workflow** so that inside **one container** you can:

1. Build/upload the Teensy firmware
2. Start the micro-ROS agent
3. Echo ROS topics

...without leaving Docker or touching the host. This will give a fully contained micro-ROS development workflow.

Do you want me to do that?

LLM request

Is there a way to put the connecting, of this node, inside the for loop and have it re-run the code if the node isn't connected (with a delay of course)

LLM response

Yes! You can definitely structure your code to **attempt to connect the micro-ROS node repeatedly** if the connection fails. The key idea is to move the `rcl_support_init` and `node/publisher` initialization into a retry loop, while leaving your main functionality (blinking LED and publishing) in `loop()`.

Here's a clean way to do it:

```
#include <Arduino.h>
#include <micro_ros_platformio.h>
#include <rcl/rcl.h>
#include <rcl/rcl.h>
#include <rcl/rcl.h>
#include <std_msgs/msg/int32.h>

#ifndef LED_BUILTIN
#define LED_BUILTIN 13
#endif

// micro-ROS objects
rcl_publisher_t publisher;
rcl_node_t node;
rcl_support_t support;
rcl_allocator_t allocator;
std_msgs__msg__Int32 msg;

bool connected = false;

void setup() {
  Serial.begin(115200);
  delay(2000);
  pinMode(LED_BUILTIN, OUTPUT);

  set_microros_serial_transports(Serial);

  allocator = rcl_get_default_allocator();

  // Attempt to connect in a loop
  while (!connected) {
    Serial.println("Trying to connect to micro-ROS agent...");
    rcl_ret_t ret;

    // Initialize support
    ret = rcl_support_init(&support, 0, NULL, &allocator);
    if (ret != RCL_RET_OK) {
      Serial.println("Failed to init support, retrying...");
      delay(1000);
      continue;
    }

    // Initialize node
    ret = rcl_node_init_default(&node, "teensy_node", "", &support);
    if (ret != RCL_RET_OK) {
      Serial.println("Failed to init node, retrying...");
      rcl_support_fini(&support);
    }
  }
}
```

```

        delay(1000);
        continue;
    }

    // Initialize publisher
    ret = rcl_publisher_init_default(
        &publisher,
        &node,
        ROSIDL_GET_MSG_TYPE_SUPPORT(std_msgs, msg, Int32),
        "teensy_counter"
    );
    if (ret != RCL_RET_OK) {
        Serial.println("Failed to init publisher, retrying...");
        rcl_node_fini(&node);
        rcl_support_fini(&support);
        delay(1000);
        continue;
    }

    // Successfully connected
    connected = true;
    Serial.println("Connected to micro-ROS agent!");
}

msg.data = 0;
}

void loop() {
    // Blink LED
    digitalWrite(LED_BUILTIN, HIGH);
    delay(500);
    digitalWrite(LED_BUILTIN, LOW);
    delay(500);

    // Publish micro-ROS message
    if (connected) {
        rcl_publish(&publisher, &msg, NULL);
        msg.data++;
    }
}

```

✓ Key Points:

1. **Retry loop in setup()** – Keeps attempting until the micro-ROS agent is reachable.
2. **Graceful cleanup** – If node or publisher fails to initialize, previous resources are cleaned with `rcl_node_fini` and `rcl_support_fini`.
3. **Delay between attempts** – Prevents spamming the agent.
4. **Flag connected** – Ensures `loop()` only tries to publish if the connection succeeded.

If you want, I can also rewrite it so that **the connection attempt happens dynamically in loop()**. That way, your program continues blinking the LED even if the agent isn't connected yet, and it automatically reconnects when the agent comes online.

Do you want me to do that version too?

LLM request

```
lib_deps = https://github.com/micro-ROS/micro_ros_platformio
```

LLM response

--

LLM request

```
lib_deps = https://github.com/micro-ROS/micro_ros_platformio
```


LLM response

--

LLM request

```
lib_deps = https://github.com/micro-ROS/micro_ros_platformio
```

LLM response

--

LLM request

Why do i get this error:

```
Why do i get this error Step 6/9 :  
RUN ~/.platformio/penv/bin/python -m pip install --upgrade pip setuptools wheel &&  
~/.platformio/penv/bin/python -m pip install catkin_pkg empy lark-parser  
---> Running in 5fc04fb8c217  
/bin/bash: line 1: /root/.platformio/penv/bin/python: No such file or directory  
The command '/bin/bash -c ~/.platformio/penv/bin/python -m pip install --upgrade pip  
setuptools wheel && ~/.platformio/penv/bin/python -m pip install catkin_pkg empy  
lark-parser' returned a non-zero code: 127
```

LLM response

Ah! I can see exactly what's happening here. The problem comes from **the order of your commands in the Dockerfile**. Let's break it down carefully.

What your Dockerfile is doing in Step 6:

```
RUN ~/.platformio/penv/bin/python -m pip install --upgrade pip
setuptools wheel && \
    ~/.platformio/penv/bin/python -m pip install catkin_pkg empy lark-
parser
```

Error:

```
/bin/bash: line 1: /root/.platformio/penv/bin/python: No such file or
directory
```

This is because **the virtual environment at ~/.platformio/penv has not yet been created**. Looking at your Dockerfile:

- Step 7 creates the venv:

```
RUN python3 -m venv ~/.platformio/penv
```

- But Step 6 tries to use ~/.platformio/penv/bin/python **before Step 7 runs**.

➔ Docker builds layers sequentially. You are referencing a path that does **not exist yet**, so bash fails.

How to fix it

You need to **create the venv first**, then install packages in it. So reorder your steps:

1) Get image

FROM ros:humble

```
SHELL ["/bin/bash", "-c"]
```

2) Install dependencies

```
RUN apt update && apt install -y \
    build-essential \
    git \
    libasio-dev \
    python3-colcon-common-extensions \
    python3-pip \
    python3-rosdep \
    python3-venv \
    && rm -rf /var/lib/apt/lists/*
```

3) Fix git issues

```
RUN git config --global pull.rebase false && \
    git config --global --add safe.directory /workspace && \
    git config --global user.name "ThomasKramp" && \
    git config --global user.email "thomas.kramp378@gmail.com"
```

4) Setup virtual environment for PlatformIO

```
RUN python3 -m venv ~/.platformio/penv && \
    ~/.platformio/penv/bin/python -m pip install --upgrade pip
setuptools wheel
```

5) Install PlatformIO and needed packages

```
RUN ~/.platformio/penv/bin/pip install platformio catkin_pkg
empy==3.3.4 lark-parser pyyaml
```

```
# 6) Change working directory
WORKDIR /workspace
```

```
# 7) Open terminal inside container
CMD ["/bin/bash"]
```

✓ Key changes:

1. **Step 4 creates the venv first.**
2. **Then Step 5 installs pip packages using the venv Python.**
3. Removed redundant source commands in RUN; they don't persist across Docker layers.

💡 Extra tip:

- Using `~/.platformio/venv/bin/python` is fine, but you can also use `python -m venv /opt/venvs/platformio` or `/root/.platformio/venv` — just be consistent.
- source commands in Docker RUN don't persist to the next layer; always call the venv Python directly.

If you want, I can rewrite your entire Dockerfile **fully corrected for PlatformIO + ROS Humble**, so it builds cleanly without errors.

Do you want me to do that?